

CS 597 REPORT 3

Yash Pradeep Gupte
MS COMPUTER SCIENCE
CWID: A20472798

1] CODE – TransGAN – Transformer GAN

(A) Image resolution = 64*64

(B) Feature extractor = Inception v3

(C) Generator_parameters = 19933952

(D) Discriminator_parameters = 2777856

(E) Train data – 1324 images

(F) Optimizers:

Adam (lr=1e-3, betas=(0.9, 0.999), eps=1e-8)

(G) Batch Size = 32

(H) Training loop

Train dataset – coco2014

Test dataset – coco2014_test

Sets a timestamp to save model and other supporting files.

Fetch image pairs ->

train_set, test_set :

rootdir -

'/Users/yashgupte21/Desktop/CS597/week2/coco_data/images16/train'

dirnames - ['class_01_zebra']

filenames – ['COCO_train2014_000000456010.png',
'COCO_train2014_000000160741.png', 'COCO_train2014_000000481760.png',

```
'COCO_train2014_000000563929.png', 'COCO_train2014_000000162181.png',  
'COCO_train2014_000000294877.png', 'COCO_train2014_000000526534.png',  
'COCO_train2014_000000090317.png', 'COCO_train2014_000000443390.png',  
'COCO_train2014_000000369736.png', .....]
```

input_data(captions16) -

```
'/Users/yashgupte21/Desktop/CS597/week2/coco_data/captions16/train/class_01_zebra/COC  
O_train2014_000000456010.txt'
```

target_data(image) -

```
'/Users/yashgupte21/Desktop/CS597/week2/coco_data/images16/train/class_01_zebra/COCO  
_train2014_000000456010.png'
```

target1_data(image_mask) -

```
'/Users/yashgupte21/Desktop/CS597/week2/coco_data/images_mask/train/class_01_zebra/C  
OCO_train2014_000000456010.png'
```

target2_data(images256) -

```
'/Users/yashgupte21/Desktop/CS597/week2/coco_data/images256/train/class_01_zebra/COC  
O_train2014_000000456010.png'
```

```
train_samples.append([input_data, target_data, target1_data,  
target2_data]) -  
[['/Users/yashgupte21/Desktop/CS597/week2/coco_data/captions16/train/class_01_zebra/CO  
CO_train2014_000000456010.txt',  
'/Users/yashgupte21/Desktop/CS597/week2/coco_data/images16/train/class_01_zebra/COCO  
_train2014_000000456010.png',  
'/Users/yashgupte21/Desktop/CS597/week2/coco_data/images_mask/train/class_01_zebra/C  
OCO_train2014_000000456010.png',  
'/Users/yashgupte21/Desktop/CS597/week2/coco_data/images256/train/class_01_zebra/COC  
O_train2014_000000456010.png'],  
['/Users/yashgupte21/Desktop/CS597/week2/coco_data/captions16/train/class_01_zebra/COC  
O_train2014_000000160741.txt',  
'/Users/yashgupte21/Desktop/CS597/week2/coco_data/images16/train/class_01_zebra/COCO  
_train2014_000000160741.png',  
'/Users/yashgupte21/Desktop/CS597/week2/coco_data/images_mask/train/class_01_zebra/C  
OCO_train2014_000000160741.png',  
'/Users/yashgupte21/Desktop/CS597/week2/coco_data/images256/train/class_01_zebra/COC  
O_train2014_000000160741.png'],...]]
```

```
train_dataset = ListDataset(train_samples, input_transform)
```

Same we will perform for test_dataset

Create_model ->

For training it will go in the else loop and create the model by setting network_data =None

Net_G ->

```
Class Generator_transformer :
```

```
    In_ch =4208
```

```
    Out_ch = 3
```

Generator_transformer architecture

```
Generator_transformer(  
    (linear): Sequential(  
        (0): Linear(in_features=768, out_features=128, bias=True)  
        (1): ReLU(inplace=True)  
    )  
    (position_embedding): Embedding(16, 128)  
    (encoder_layers): TransformerEncoderLayer(  
        (self_attn): MultiheadAttention(  
            (out_proj): NonDynamicallyQuantizableLinear(in_features=128, out_features=128,  
bias=True)  
        )  
        (linear1): Linear(in_features=128, out_features=128, bias=True)  
        (dropout): Dropout(p=0.2, inplace=False)  
        (linear2): Linear(in_features=128, out_features=128, bias=True)  
        (norm1): LayerNorm((128,), eps=1e-05, elementwise_affine=True)  
        (norm2): LayerNorm((128,), eps=1e-05, elementwise_affine=True)  
        (dropout1): Dropout(p=0.2, inplace=False)  
        (dropout2): Dropout(p=0.2, inplace=False)  
    )  
    (transformer_encoder): TransformerEncoder(  
        (layers): ModuleList(  
            (0): TransformerEncoderLayer(  
                (self_attn): MultiheadAttention(  
                    (out_proj): NonDynamicallyQuantizableLinear(in_features=128, out_features=128,  
bias=True)  
                )  
                (linear1): Linear(in_features=128, out_features=128, bias=True)  
                (dropout): Dropout(p=0.2, inplace=False)
```

```

        (linear2): Linear(in_features=128, out_features=128, bias=True)
        (norm1): LayerNorm((128,), eps=1e-05, elementwise_affine=True)
        (norm2): LayerNorm((128,), eps=1e-05, elementwise_affine=True)
        (dropout1): Dropout(p=0.2, inplace=False)
        (dropout2): Dropout(p=0.2, inplace=False)
    )
    (1): TransformerEncoderLayer(
      (self_attn): MultiheadAttention(
        (out_proj): NonDynamicallyQuantizableLinear(in_features=128, out_features=128,
bias=True)
      )
      (linear1): Linear(in_features=128, out_features=128, bias=True)
      (dropout): Dropout(p=0.2, inplace=False)
      (linear2): Linear(in_features=128, out_features=128, bias=True)
      (norm1): LayerNorm((128,), eps=1e-05, elementwise_affine=True)
      (norm2): LayerNorm((128,), eps=1e-05, elementwise_affine=True)
      (dropout1): Dropout(p=0.2, inplace=False)
      (dropout2): Dropout(p=0.2, inplace=False)
    )
  )
)
(deconv0): Sequential(
  (0): ConvTranspose2d(2048, 512, kernel_size=(4, 4), stride=(2, 2), bias=False)
  (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (2): ReLU(inplace=True)
)
(deconv1): Sequential(
  (0): ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1),
bias=False)
  (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (2): ReLU(inplace=True)
)
(deconv2): Sequential(
  (0): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1),
bias=False)
  (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (2): ReLU(inplace=True)
)
(deconv3): Sequential(
  (0): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1),
bias=False)

```

```

    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): ReLU(inplace=True)
  )
  (deconv4): outdeconv1(
    (conv): Sequential(
      (0): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1),
bias=False)
      (1): Tanh()
    )
  )
)
)
)
)

```

```

M = BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

Net_d →

Class Discriminator :

```

Initializer - d=64, h=256, w=256
Nc = 3
Ndf = 64

```

Discriminator architecture

```

Sequential(
  (0): Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
  (1): LeakyReLU(negative_slope=0.2, inplace=True)
  (2): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
  (3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (4): LeakyReLU(negative_slope=0.2, inplace=True)
  (5): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
  (6): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (7): LeakyReLU(negative_slope=0.2, inplace=True)
  (8): Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
  (9): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (10): LeakyReLU(negative_slope=0.2, inplace=True)
)

```

Losses ->

Epoch_size = 42

(1) Update D network: maximize $D(x) - D(G(z))$

#train with real

```
netd.zero_grad()
errD_real = netd(real_var, cap_vectors_var).mean(0).view(1)
errD_real.backward(one)
```

train with fake

```
fake = netg(cap_vectors_var, mask_var)
errD_fake = netd(fake.detach(), cap_vectors_var).mean(0).view(1)
errD_fake.backward(mone)
```

total discriminator loss

```
errD = errD_real - errD_fake
netd_losses.update(errD.mean().view(1).data, batch_size)
netd_optimizer.step()

netd_writer.add_scalar('net_err_minibatch', errD.mean().view(1), n_iter)
```

(2) Update G network: maximize $D(G(z))$

```
netg.zero_grad()
fake = netg(cap_vectors_var, mask_var)

feature_real = featureExtractorNet(real_var)
feature_real = feature_real.detach()
```

```
feature_fake = featureExtractorNet(fake)
feature_fake = feature_fake.detach()
```

```
adv_loss = netd(fake, cap_vectors_var).mean(0).view(1)
content_loss = contentLoss(fake, real_var)
```

```
errG = adv_loss + 100.0 * content_loss
```

```
errG.backward(one)
```

```
netg_losses.update(errG.mean().view(1).data, batch_size)
netg_optimizer.step()
```

```
netg_writer.add_scalar('net_err_minibatch', errG.mean().view(1), n_iter)
```

(I) Running time:

It took approximately 5 minutes to run one epoch on my machine